

Copyright 2020, MIT Lincoln Laboratory

SPDX-License-Identifier: BSD-2-Clause

▼ 1 Example of Moments in Time Model Inference

The following loads and parses a video from the "Skiing" class of the UCF101 dataset. This notebook shows a minimum example of inference on this video which the model has not previously seen.

View and (optionally) modify the TODOs below to try your own video on a selection of these models.

▼ 1.1 Setup

▼ 1.2 Get UCF101 video dataset

Go to <https://www.crcv.ucf.edu/data/UCF101.php> (<https://www.crcv.ucf.edu/data/UCF101.php>) and download the dataset. Locate the 'Skiing' class and choose a video.

▼ 1.2.1 Import Packages and Methods

```
In [1]: 1 import os
        2 import numpy as np
        3 import re
        4 import sys
        5 import glob
        6 import argparse
        7 import functools
        8 import subprocess
        9 from PIL import Image
       10 import time
       11 import torch
       12 from torchvision import transforms as trn
       13 import h5py
       14 import shutil
       15 from tensorflow.python.keras.utils import to_categorical
       16
```

```
/state/partition1/llgrid/pkg/anaconda/anaconda3-2020a/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:516: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,) type'.
```

```
_np_qint8 = np.dtype [("qint8", np.int8, 1)]
/state/partition1/llgrid/pkg/anaconda/anaconda3-2020a/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:517: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,) type'.
```

```
_np_quint8 = np.dtype [("quint8", np.uint8, 1)]
/state/partition1/llgrid/pkg/anaconda/anaconda3-2020a/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:518: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,) type'.
```

```
_np_qint16 = np.dtype [("qint16", np.int16, 1)]
/state/partition1/llgrid/pkg/anaconda/anaconda3-2020a/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:519: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,) type'.
```

```
_np_quint16 = np.dtype [("quint16", np.uint16, 1)]
/state/partition1/llgrid/pkg/anaconda/anaconda3-2020a/lib/python3.6/site-packages/tensorflow/python/framework/dtypes.py:520: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,) type'.
```

```
WARNING:tensorflow:From /state/partition1/llgrid/pkg/anaconda/anaconda3-2020a/lib/python3.6/site-packages/horovod/tensorflow/__init__.py:117: The name tf.global_variables is deprecated. Please use tf.compat.v1.global_variables instead.
```

```
WARNING:tensorflow:From /state/partition1/llgrid/pkg/anaconda/anaconda3-2020a/lib/python3.6/site-packages/horovod/tensorflow/__init__.py:143: The name tf.get_default_graph is deprecated. Please use tf.compat.v1.get_default_graph instead.
```

```
/state/partition1/llgrid/pkg/anaconda/anaconda3-2020a/lib/python3.6/site-packages/tensorboard/compat/tensorflow_stub/dtypes.py:541: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,) type'.
```

```
_np_qint8 = np.dtype [("qint8", np.int8, 1)]
/state/partition1/llgrid/pkg/anaconda/anaconda3-2020a/lib/python3.6/site-packages/tensorboard/compat/tensorflow_stub/dtypes.py:542: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,)) / '(1,) type'.
```

```
_np_quint8 = np.dtype [("quint8", np.uint8, 1)]
/state/partition1/llgrid/pkg/anaconda/anaconda3-2020a/lib/python3.6/site-packages/tensorboard/compat/tensorflow_stub/dtypes.py:543: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,))
```

```

/ '(1,)type'.
_np_qint16 = np.dtype [("qint16", np.int16, 1)]
/state/partition1/llgrid/pkg/anaconda/anaconda3-2020a/lib/python3.6/site-packages/tensorboard/compat/tensorflow_stub/dtypes.py:544: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,))
/ '(1,)type'.
_np_quint16 = np.dtype [("quint16", np.uint16, 1)]
/state/partition1/llgrid/pkg/anaconda/anaconda3-2020a/lib/python3.6/site-packages/tensorboard/compat/tensorflow_stub/dtypes.py:545: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,))
/ '(1,)type'.
_np_qint32 = np.dtype [("qint32", np.int32, 1)]
/state/partition1/llgrid/pkg/anaconda/anaconda3-2020a/lib/python3.6/site-packages/tensorboard/compat/tensorflow_stub/dtypes.py:550: FutureWarning: Passing (type, 1) or 'ltype' as a synonym of type is deprecated; in a future version of numpy, it will be understood as (type, (1,))
/ '(1,)type'.
np_resource = np.dtype [("resource", np.ubyte, 1)]

```

▼ 1.2.2 Enable Model File Reading

```
In [2]: 1 os.putenv("HDF5_USE_FILE_LOCKING", "FALSE")
```

```
Out[2]: 0
```

▼ 1.2.3 Create Moments in Time One-Hot Class Labels

```
In [3]: 1 set_dir = '/home/gridsan/groups/Moments_in_Time/data-copy/data/parsed/
2 # Read folders in set directory. Folders should correspond to classes
3 categories = os.listdir(set_dir)
4 # Remove any extraneous .ipynb files
5 new_categories = []
6 for cat in categories:
7     if '.ipynb' not in cat:
8         new_categories.append(cat)
9 categories = new_categories
10 # Create one-hot vector labels
11 categories_labels = dict()
12 for i in range(len(categories)):
13     categories_labels[categories[i]] = to_categorical(np.array(i), num
14     ...
339 total classes

```

▼ 1.3 Parse Video

▼ 1.3.1 Specify Input

```
In [4]: 1 # TODO:
```

```

2 # Specify input video and corresponding ground truth Moments in Time d
3 video_file = 'example_video.avi' ### TODO: add your video file from th
4 ucf101_class_gt = 'Skiing'

```

▼ 1.3.2 Load and Parse

```

In [5]: 1 def read_as_list(filename):
2         """Read the file at filename and store its contents into a list."""
3         with open(filename) as f:
4             result = [line.rstrip() for line in f.readlines()]
5             f.close()
6             return result
7
8 def load_frames(frame_paths, num_frames):
9     """Loads frames from file and returns a list of the frames."""
10    #print('Running load_frames')
11    frames = [Image.open(frame).convert('RGB') for frame in frame_paths]
12    if len(frames) >= num_frames:
13        return frames[:int(np.ceil(len(frames) / float(num_frames)))]
14    else:
15        raise ValueError('Video must have at least {} frames'.format(num_frames))
16
17 def extract_frames(video_file, framesFolder, framerate):
18     """Takes a video and converts it into a list of frames."""
19     if os.path.exists(framesFolder):
20         subprocess.call(['rm', '-rf', framesFolder + '*.jpg'])
21     try:
22         os.makedirs(os.path.join(framesFolder))
23     except OSError:
24         pass
25     output = subprocess.Popen(['ffmpeg', '-i', video_file], stderr=subprocess.PIPE)
26     # Search and parse 'Duration: 00:05:24.13,' from ffmpeg stderr
27     re_duration = re.compile('Duration: (.*?)\.')
28     duration = re_duration.search(str(output.stderr.readlines()[0])).groups()[0]
29     seconds = functools.reduce(lambda x, y: x * 60 + y,
30                               map(int, duration.split(':')))
31     #rate = num_frames / float(seconds)
32     rate = framerate
33     num_frames = seconds * rate
34     output = subprocess.Popen(['ffmpeg', '-i', video_file,
35                               '-vf', 'fps={}'.format(rate),
36                               '-vframes', str(num_frames),
37                               '-loglevel', 'panic',
38                               os.path.join(framesFolder, '%d.jpg')] +
39                               [os.path.join(framesFolder, frame)
40                               for frame in os.listdir(framesFolder)])
41     frames = load_frames([os.path.join(framesFolder, frame)
42                           for frame in os.listdir(framesFolder)], num_frames)
43     subprocess.call(['rm', '-rf', framesFolder + '*.jpg'])
44     return frames
45
46 def remove(path):
47     """ param <path> could either be relative or absolute. """
48     if os.path.isfile(path) or os.path.islink(path):
49         os.remove(path) # remove the file
50     elif os.path.isdir(path):

```

```

50         shutil.rmtree(path) # remove dir and all contains
51     else:
52         raise ValueError("file {} is not a file or dir.".format(path))
53
54 def parse(video_file):
55     """Extracts a numpy array representation from a given video file."
56     # Extract frames from video
57     framerate = 25 # frames per second (fps)
58     framesFolder = os.getcwd() + '/frames'
59     frames = extract_frames(video_file, framesFolder, framerate)
60     # Transform frames into tensors with values [0,1]
61     # Load an image transformer
62     transform = trn.ToTensor()
63     tframes = []
64     for frame in frames:
65         #frame = frame.resize((224,224))
66         tframes.append(frame)
67     tmpdata = torch.stack([transform(frame) for frame in tframes] )
68     tmpdata = tmpdata.numpy()
69     remove(framesFolder)
70     return tmpdata
71
72 video_array = parse(video_file)
Video shape: (200, 3, 240, 320)

```

▼ 1.4 Prepare Model(s)

▼ 1.4.1 Select Trained Model(s)

```

In [6]: 1 # TODO:
2 # Select models from: ['DenseNet169', 'DenseNet201', 'InceptionV3', 'I
3 #                               'MobileNetV2', 'Resnet50', 'VGG19', 'Xception',
4 #                               'InceptionResnetV2-64avg']
5
6 model_names = ['DenseNet169', 'DenseNet201', 'InceptionV3', 'Inception
7                 'MobileNetV2', 'Resnet50', 'VGG19', 'Xception', 'I3D-In
8                 'InceptionResNetV2-64avg']
9
10 # TODO:
11 # Specify model directory
12

```

▼ 1.4.2 Load and Compile Model(s)

```

In [7]: 1 def get_model(model_name):
2         if model_name == 'DenseNet169':
3             return load_model(model_folder + 'D169-224x224x3-339-im.h5')
4         elif model_name == 'DenseNet201':
5             return load_model(model_folder + 'D201-224x224x3-339-im.h5')
6         elif model_name == 'InceptionV3':
7             return load_model(model_folder + 'Iv3-224x224x3-339-im.h5')
8         elif model_name == 'InceptionResNetV2':

```

```

9         return load_model(model_folder + 'IRv2-224x224x3-339-im.h5')
10    elif model_name == 'MobileNet':
11        return load_model(model_folder + 'M-224x224x3-339-im.h5')
12    elif model_name == 'MobileNetV2':
13        return load_model(model_folder + 'Mv2-224x224x3-339-im.h5')
14    elif model_name == 'Resnet50':
15        return load_model(model_folder + 'R50-224x224x3-339-im.h5')
16    elif model_name == 'VGG19':
17        return load_model(model_folder + 'VGG19-224x224x3-339-im.h5')
18    elif model_name == 'Xception':
19        return load_model(model_folder + 'X-224x224x3-339-im.h5')
20    elif model_name == 'C3D':
21        return load_model(model_folder + 'C3D-16x224x224x3-339-m.h5')
22    elif model_name == 'I3D-InceptionV1':
23        return load_model(model_folder + 'I3DInv1-16x224x224x3-339-ikm.')
24    elif model_name == 'LRCN':
25        return load_model(model_folder + 'LRCN-16x224x224x3-339-m.h5')
26    elif model_name == 'InceptionResNetV2-64avg':
27        return load_model(model_folder + 'IRv2avg-64x224x224x3-339-im.')
28    print('Should not reach here')
29    return None
30
31    def transformer(model_name):
32        if model_name in ['DenseNet169', 'DenseNet201', 'InceptionV3', 'In
33            'MobileNetV2', 'Resnet50', 'VGG19', 'Xception']:
34            return 1
35        elif model_name in ['I3D-InceptionV1', 'C3D', 'LRCN']:
36            return 2
37        elif model_name in ['InceptionResNetV2-64avg']:
38            return 3
39        else:
40            print('Should not reach here')
41            return None
42
43    met = get_metrics()
44    models = [get_model(m) for m in model_names]
45    transforms = [transformer(m) for m in model_names]
46    for model in models:
47        model.compile(optimizer=SGD(), loss='categorical_crossentropy', me
48    print('-----')

```

WARNING:tensorflow:From /state/partition1/llgrid/pkg/anaconda/anaconda3-2020a/lib/python3.6/site-packages/tensorflow/python/ops/init_ops.py:97: calling GlorotUniform.__init__ (from tensorflow.python.ops.init_ops) with dtype is deprecated and will be removed in a future version.

Instructions for updating:

Call initializer instance with the dtype argument instead of passing it to the constructor

WARNING:tensorflow:From /state/partition1/llgrid/pkg/anaconda/anaconda3-2020a/lib/python3.6/site-packages/tensorflow/python/ops/init_ops.py:1251: calling VarianceScaling.__init__ (from tensorflow.python.ops.init_ops) with dtype is deprecated and will be removed in a future version.

Instructions for updating:

Call initializer instance with the dtype argument instead of passing it to the constructor

WARNING:tensorflow:From /state/partition1/llgrid/pkg/anaconda/anaconda

```

3-2020a/lib/python3.6/site-packages/tensorflow/python/ops/init_ops.py:
97: calling Zeros.__init__ (from tensorflow.python.ops.init_ops) with
dtype is deprecated and will be removed in a future version.
Instructions for updating:
Call initializer instance with the dtype argument instead of passing i
t to the constructor
WARNING:tensorflow:From /state/partition1/llgrid/pkg/anaconda/anaconda
3-2020a/lib/python3.6/site-packages/tensorflow/python/ops/init_ops.py:
97: calling Ones.__init__ (from tensorflow.python.ops.init_ops) with d
type is deprecated and will be removed in a future version.
Instructions for updating:
Call initializer instance with the dtype argument instead of passing i
t to the constructor
WARNING:tensorflow:From /state/partition1/llgrid/pkg/anaconda/anaconda
3-2020a/lib/python3.6/site-packages/tensorflow/python/ops/init_ops.py:
97: calling Orthogonal.__init__ (from tensorflow.python.ops.init_ops)
with dtype is deprecated and will be removed in a future version.
Instructions for updating:
Call initializer instance with the dtype argument instead of passing i
t to the constructor
WARNING:tensorflow:From /state/partition1/llgrid/pkg/anaconda/anaconda
3-2020a/lib/python3.6/site-packages/tensorflow/python/ops/math_grad.p
y:1250: add_dispatch_support.<locals>.wrapper (from tensorflow.python.
ops.array_ops) is deprecated and will be removed in a future version.
Instructions for updating:
Use tf.where in 2.0, which has the same broadcast rule as np.where
WARNING:tensorflow:No training configuration found in save file: the m
odel was *not* compiled. Compile it manually.
-----
13 model(s) loaded and compiled

```



1.5 Inference

Note that there is still randomness in the frame(s) being selected for inference. Therefore, running this cell multiple times will give slightly different results.

```

In [8]: 1 # TODO:
        2 # Specify number of top prediction classes to show
        3

```

```

In [9]: 1 # Performs Inference and Displays Results
        2
        3 def label_to_cat(index, classes):
        4     for cat in categories_labels.keys():
        5         if categories_labels[cat][index]==1:
        6             return cat
        7 def top_classes(tups, k=5):
        8     top_k = []
        9     for i in range(len(tups)):
       10         if tups[i][1] <= k:
       11             top_k.append(tups[i])
       12     return sorted(top_k, key=lambda tup: tup[1])
       13
       14
       15 label = categories_labels[mit_class_gt]

```

```

16 datum = [video_array], 0, label
17
18 print('Model Name')
19 print('(class_name, prediction_rank, probability)')
20 print('-----')
21 print()
22
23 for i in range(len(models)):
24     model_name = model_names[i]
25     model = models[i]
26     t = transforms[i]
27     y_hats = []
28     for i in range(5): # sample averaging
29         arr, label = transform(datum, transform_num=t)
30         X = np.array([arr])
31         y = np.array([label])
32         y_hat = model.predict(X, batch_size=1)
33         y_hats.append(y_hat)
34     y_hat = np.average(y_hats, axis=0)
35     order = y_hat.argsort()
36     rev_ranks = order.argsort()
37     classes = len(y_hat[0])
38     tups = []
39     for i in range(classes):
40         cat = label_to_cat(i, classes)
41         tups.append( (cat, 339-rev_ranks[0][i], y_hat[0][i]) )
42     short_list = top_classes(tups, k)
43     print(model_name)
44     for tup in short_list:
45         print(tup)

```

```

Model Name
(class_name, prediction_rank, probability)
-----

```

```

DenseNet169
('skiing', 1, 0.9085125)
('boarding', 2, 0.025353933)
('sliding', 3, 0.0076259063)
('slipping', 4, 0.005177059)
('skating', 5, 0.0050613997)

```

```

DenseNet201
('skiing', 1, 0.58287513)
('sliding', 2, 0.05612815)
('slipping', 3, 0.043456513)
('skating', 4, 0.03366632)
('descending', 5, 0.018783813)

```

```

InceptionV3
('skiing', 1, 0.5834716)
('sliding', 2, 0.028677518)
('slipping', 3, 0.028165314)
('skating', 4, 0.023576682)
('boarding', 5, 0.021372344)

```


InceptionResNetV2

```
('skiing', 1, 0.6004424)
('boarding', 2, 0.053444456)
('jumping', 3, 0.034714162)
('skating', 4, 0.02377763)
('falling', 5, 0.023125973)
```

MobileNet

```
('skiing', 1, 0.2398886)
('officiating', 2, 0.14038408)
('slipping', 3, 0.10001288)
('boarding', 4, 0.095631294)
('sliding', 5, 0.038019247)
```

MobileNetV2

```
('skiing', 1, 0.16444364)
('sliding', 2, 0.07168311)
('boarding', 3, 0.06656089)
('slipping', 4, 0.06320612)
('skating', 5, 0.036344867)
```

Resnet50

```
('skiing', 1, 0.6975399)
('slipping', 2, 0.035883404)
('skating', 3, 0.023647005)
('balancing', 4, 0.015892975)
('officiating', 5, 0.015808614)
```

VGG19

```
('slipping', 1, 0.113963984)
('boarding', 2, 0.105429515)
('skiing', 3, 0.0875464)
('sliding', 4, 0.06347631)
('skating', 5, 0.043519128)
```

Xception

```
('skiing', 1, 0.45929837)
('slipping', 2, 0.13140288)
('racing', 3, 0.03646811)
('boarding', 4, 0.02920106)
('playing+sports', 5, 0.027381152)
```

I3D-InceptionV1

```
('skiing', 1, 0.6306972)
('boarding', 2, 0.04456868)
('sliding', 3, 0.03136454)
('skating', 4, 0.029156972)
('playing+sports', 5, 0.028053453)
```

C3D

```
('skiing', 1, 0.1877743)
('skating', 2, 0.07308701)
('crashing', 3, 0.032002218)
('boarding', 4, 0.02330643)
('flipping', 5, 0.02312484)
```

```
LRCN
('skiing', 1, 0.3637656)
('skating', 2, 0.07882427)
('slipping', 3, 0.041725338)
('jumping', 4, 0.03955111)
('sliding', 5, 0.035640597)
```

```
InceptionResNetV2-64avg
('skiing', 1, 0.57533085)
('boarding', 2, 0.074289046)
('slipping', 3, 0.035638716)
('jumping', 4, 0.028740693)
('sliding', 5, 0.019542733)
```



1.6 Questions

Any questions can be directed to Matthew Hutchinson at hutchinson@alum.mit.edu (<mailto:hutchinson@alum.mit.edu>).

Python license: <https://docs.python.org/3/license.html> (<https://docs.python.org/3/license.html>)

TensorFlow license: <https://github.com/tensorflow/tensorflow/blob/master/LICENSE> (<https://github.com/tensorflow/tensorflow/blob/master/LICENSE>)

NumPy license: <https://numpy.org/doc/stable/license.html> (<https://numpy.org/doc/stable/license.html>)

PIL license: <http://www.pythonware.com/products/pil/license.htm> (<http://www.pythonware.com/products/pil/license.htm>)

PyTorch license: <https://github.com/pytorch/pytorch/blob/master/LICENSE> (<https://github.com/pytorch/pytorch/blob/master/LICENSE>)

H5Py license: <https://docs.h5py.org/en/stable/licenses.html> (<https://docs.h5py.org/en/stable/licenses.html>)

In []:



Present



Slides



Themes



Help